

# Stream processing - part III (real-time data architecture)

Bartosz Łoś  
RTB House

# What will this lecture be about?

Goal:

- show an example of real-time data architecture at RTB House (different approaches and use cases, design decisions)
- dig deep into data processing frameworks (Apache Storm, Kafka Streams, Kafka Workers etc.)



# The context: RTB platform


Our platform:

- takes part in auctions, purchases and emits advertisements in the RTB model
- processes **5M bid requests per second** and generates **300K events per second** (**200TB data every day**)

Data processing:

- requirements: machine learning, system monitoring (alerting, ad hoc debugging) and financial settlements (reports, budget limits)
- use cases: filtering, synchronizing, joining, aggregating, storing events and statistics in Hadoop, BigQuery, Postgres or Elasticsearch

# The context: RTB platform



RTB HOUSE

# The context: RTB platform

Our platform consists of two types of servlets:

- **bidders** process **bid requests**
- **adservlets** process **user requests** (tags, impressions, clicks and conversions)



# The context: RTB platform

To be able to buy advertising space effectively, we needed to store and process data (user info, historical impressions)

We were able to use this data for estimating:

- probability of a click (click-through rate) - **CTR**
- conditional probability of a conversion given that an impression was clicked (conversion rate) - **CR**
- conversion value - **CV**

These estimated values are used for bid pricing:

$$\text{bid\_value} = (1 - \text{margin}) * \mathbf{CTR} * \mathbf{CR} * \mathbf{CV} * \text{rate}$$

# Iterations

We have been improving our solution by many iterations:

- at first: end-of-day batch jobs, single-DC, inconsistent data-flows
- finally: real-time data processing, delay reduced from 1 day to 15 seconds, multi-DC architecture, end-to-end exactly-once processing

It was essential to:

- separate data-flow from the core platform
- provide immutable streams of events and data synchronization between DCs
- dig deep into open-source streaming technologies and if needed replace them by better, custom-built components

# The 1st iteration: mutable impressions

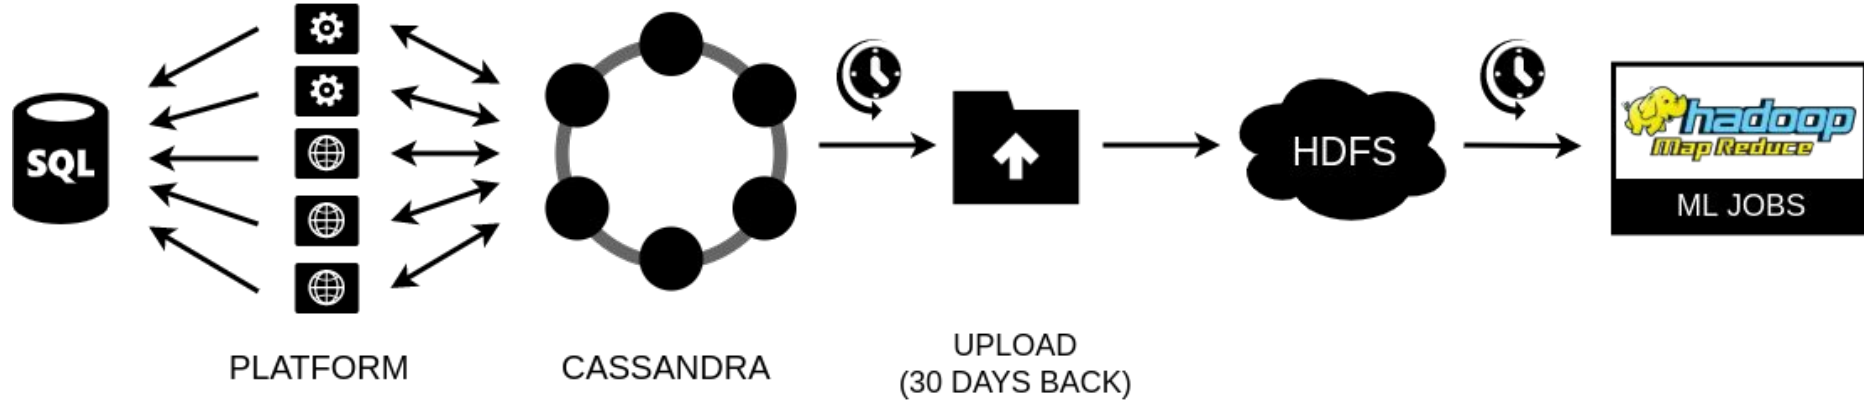
```
{ IMPRESSION:  
  IMPRESSION_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  ...  
  CLICKS,  
  CONVERSIONS  
}
```

```
{ CLICK:  
  CLICK_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  IMPRESSION_HASH,  
  ...  
}
```

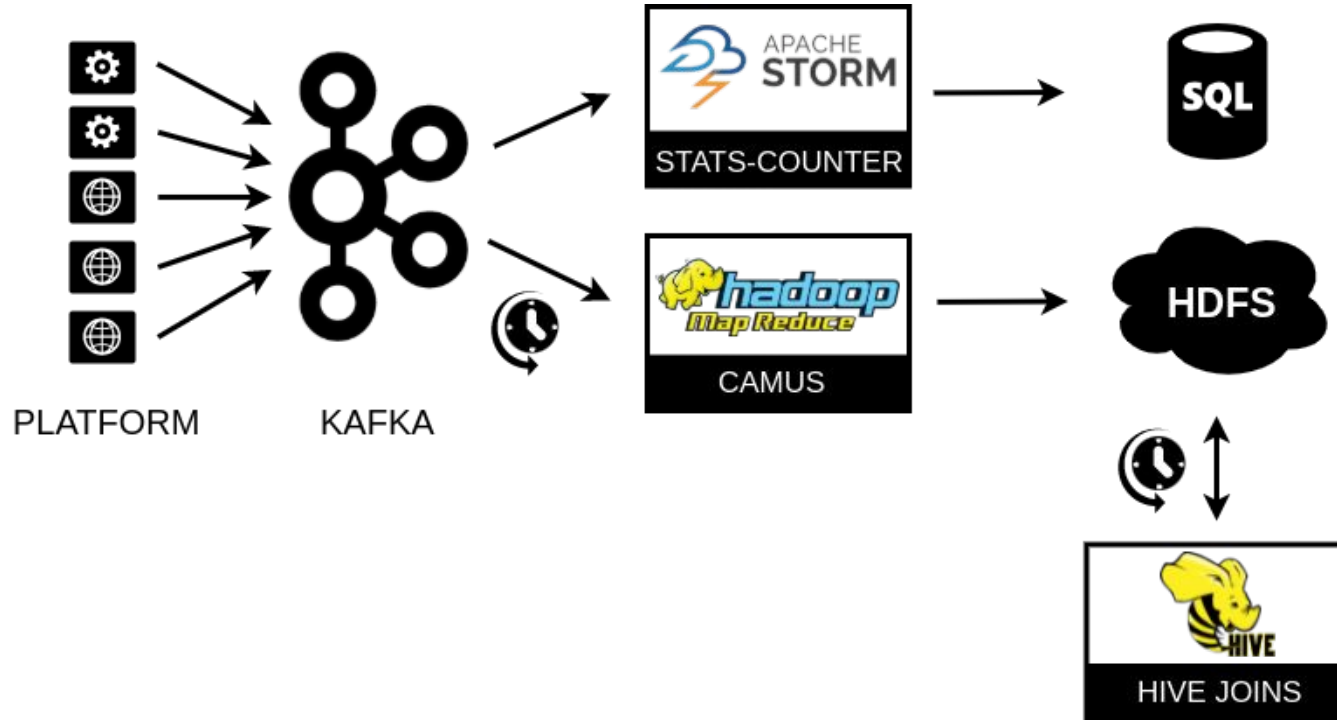
```
{ CONVERSION:  
  CONVERSION_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  ...  
}
```



# The 1st iteration: mutable impressions



# The 2nd iteration: data-flow



# The 3rd iteration: immutable streams of events

```
{ IMPRESSION:  
  IMPRESSION_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  ...  
  CLICKS,  
  CONVERSIONS  
}
```

```
{ CLICK:  
  CLICK_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  IMPRESSION_HASH,  
  ...  
}
```

```
{ CONVERSION:  
  CONVERSION_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  ...  
}
```

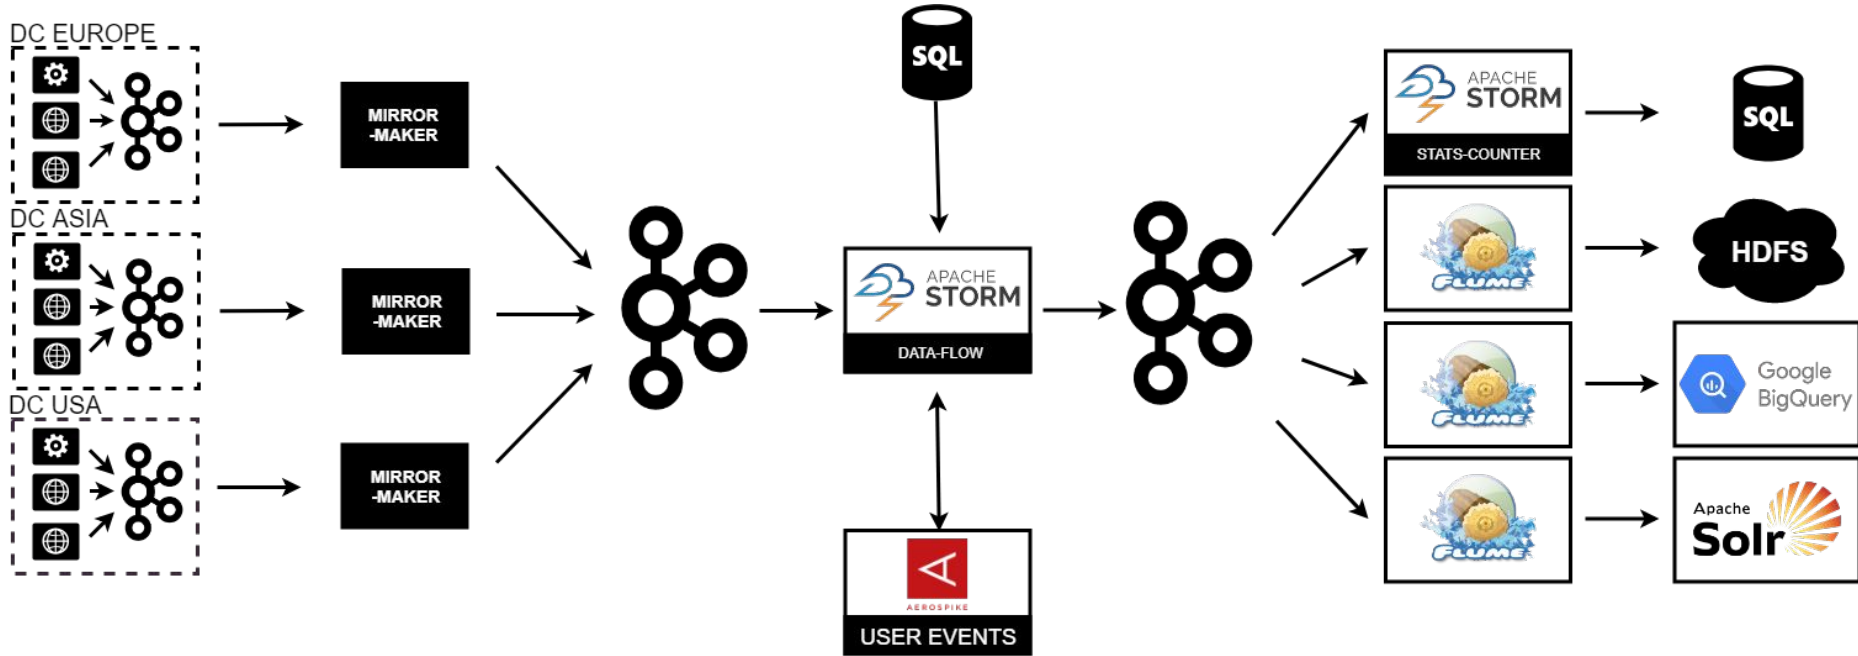
# The 3rd iteration: immutable streams of events

```
{ IMPRESSION:  
  IMPRESSION_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  ...  
  CLICKS,  
  CONVERSIONS  
}
```

```
{ CLICK:  
  CLICK_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  IMPRESSION_HASH,  
  ...  
  IMPRESSION  
}
```

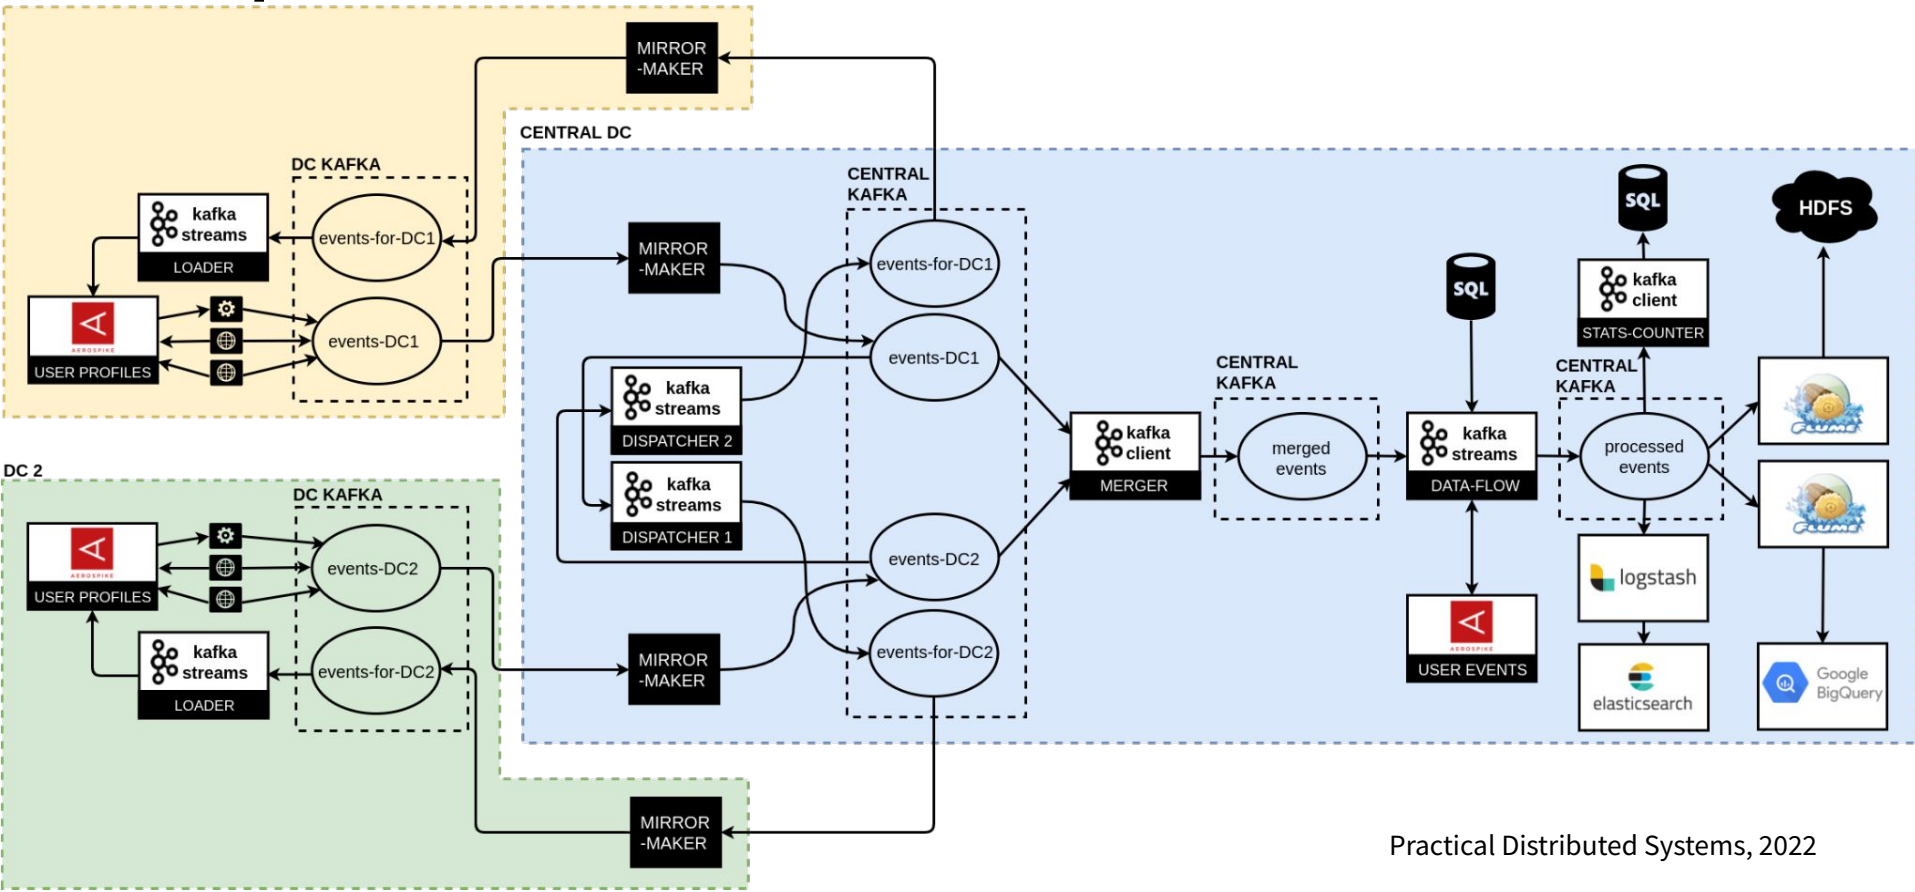
```
{ CONVERSION:  
  CONVERSION_HASH,  
  TIME,  
  COOKIE,  
  ADVERTISER_ID,  
  ...  
  IMPRESSION,  
  CLICK  
}
```

# The 3rd iteration: immutable streams of events



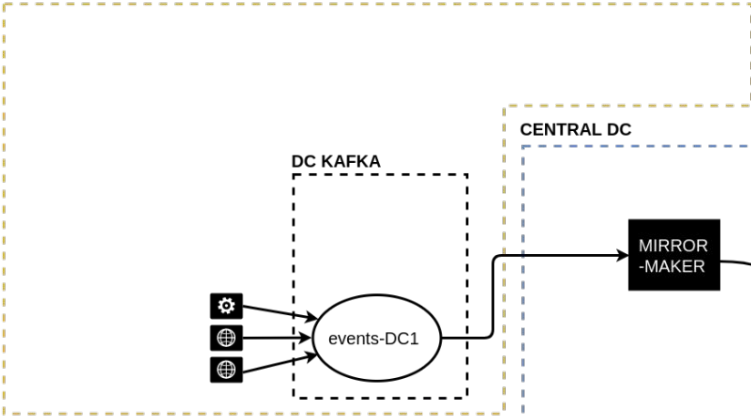
# The 4th iteration: multi-dc architecture

DC 1

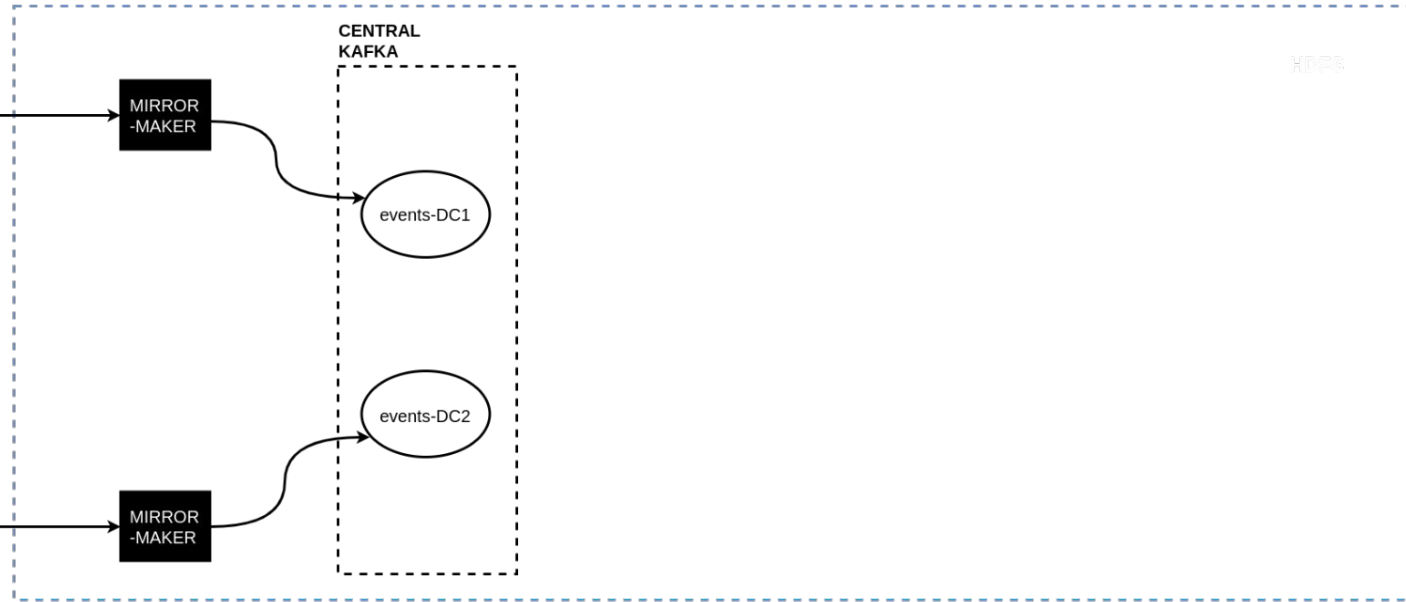


# The 4th iteration: multi-dc architecture

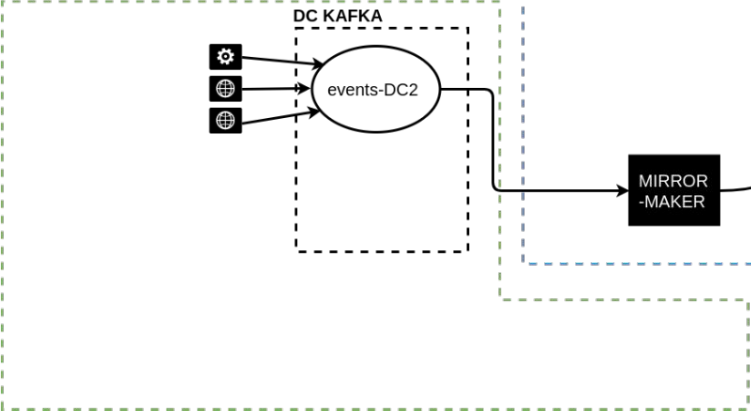
DC 1



CENTRAL DC

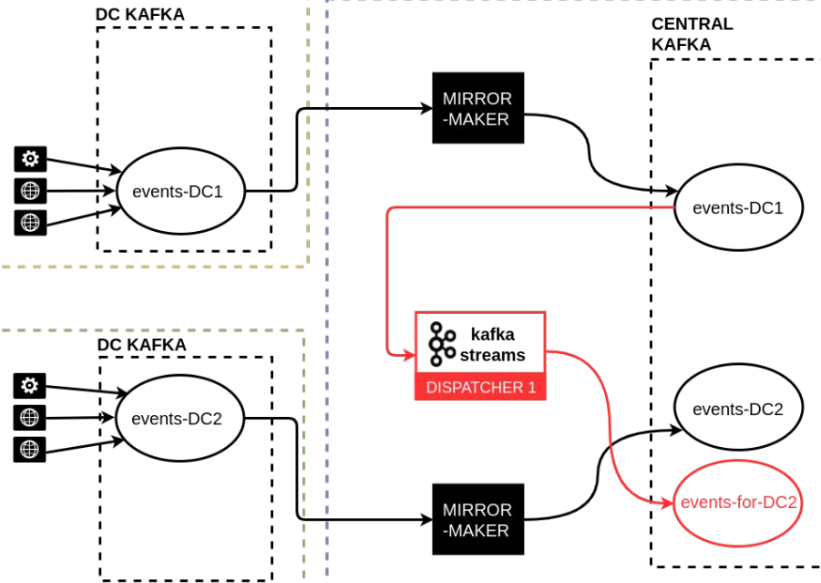


DC 2



# The 4th iteration: multi-dc architecture

DC 1

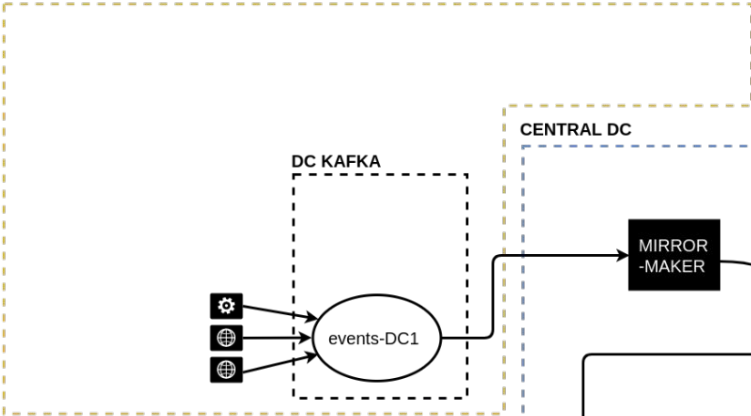


KDCS

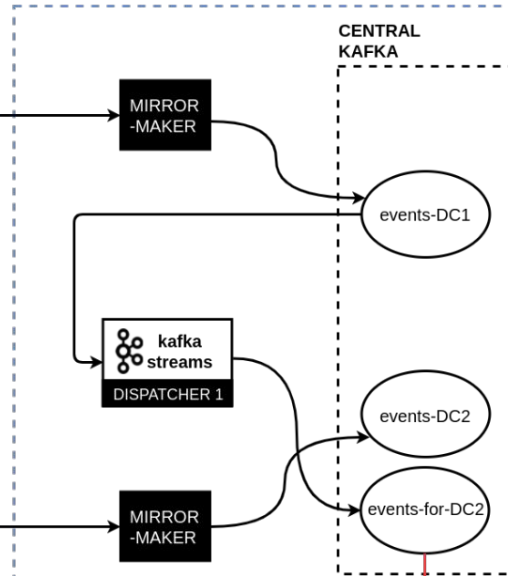


# The 4th iteration: multi-dc architecture


DC 1



CENTRAL DC



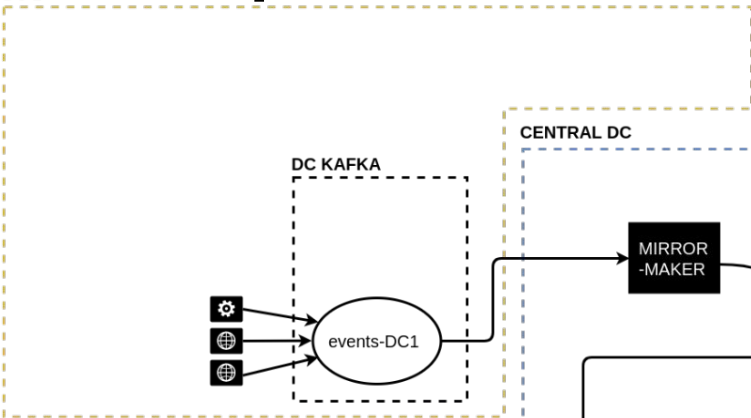
DC 2



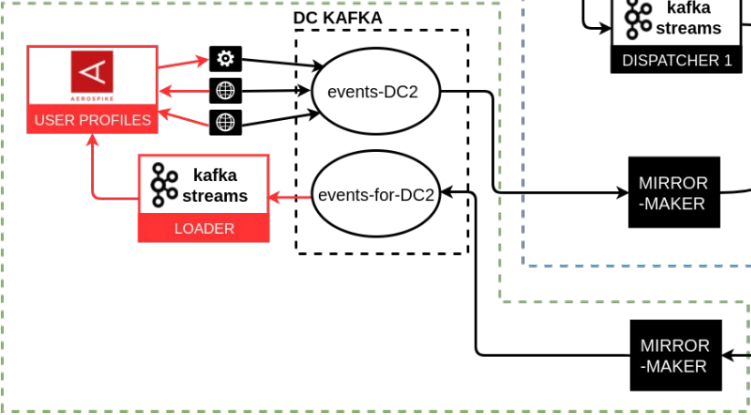
MIRROR-MAKER

# The 4th iteration: multi-dc architecture

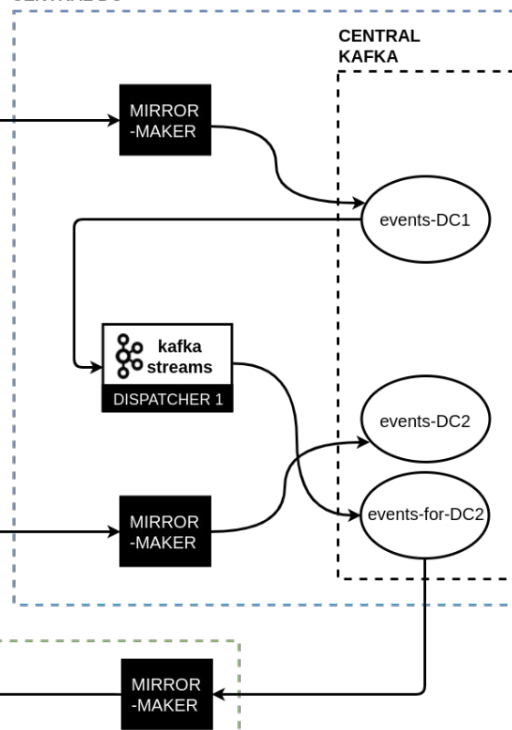
DC 1



DC 2

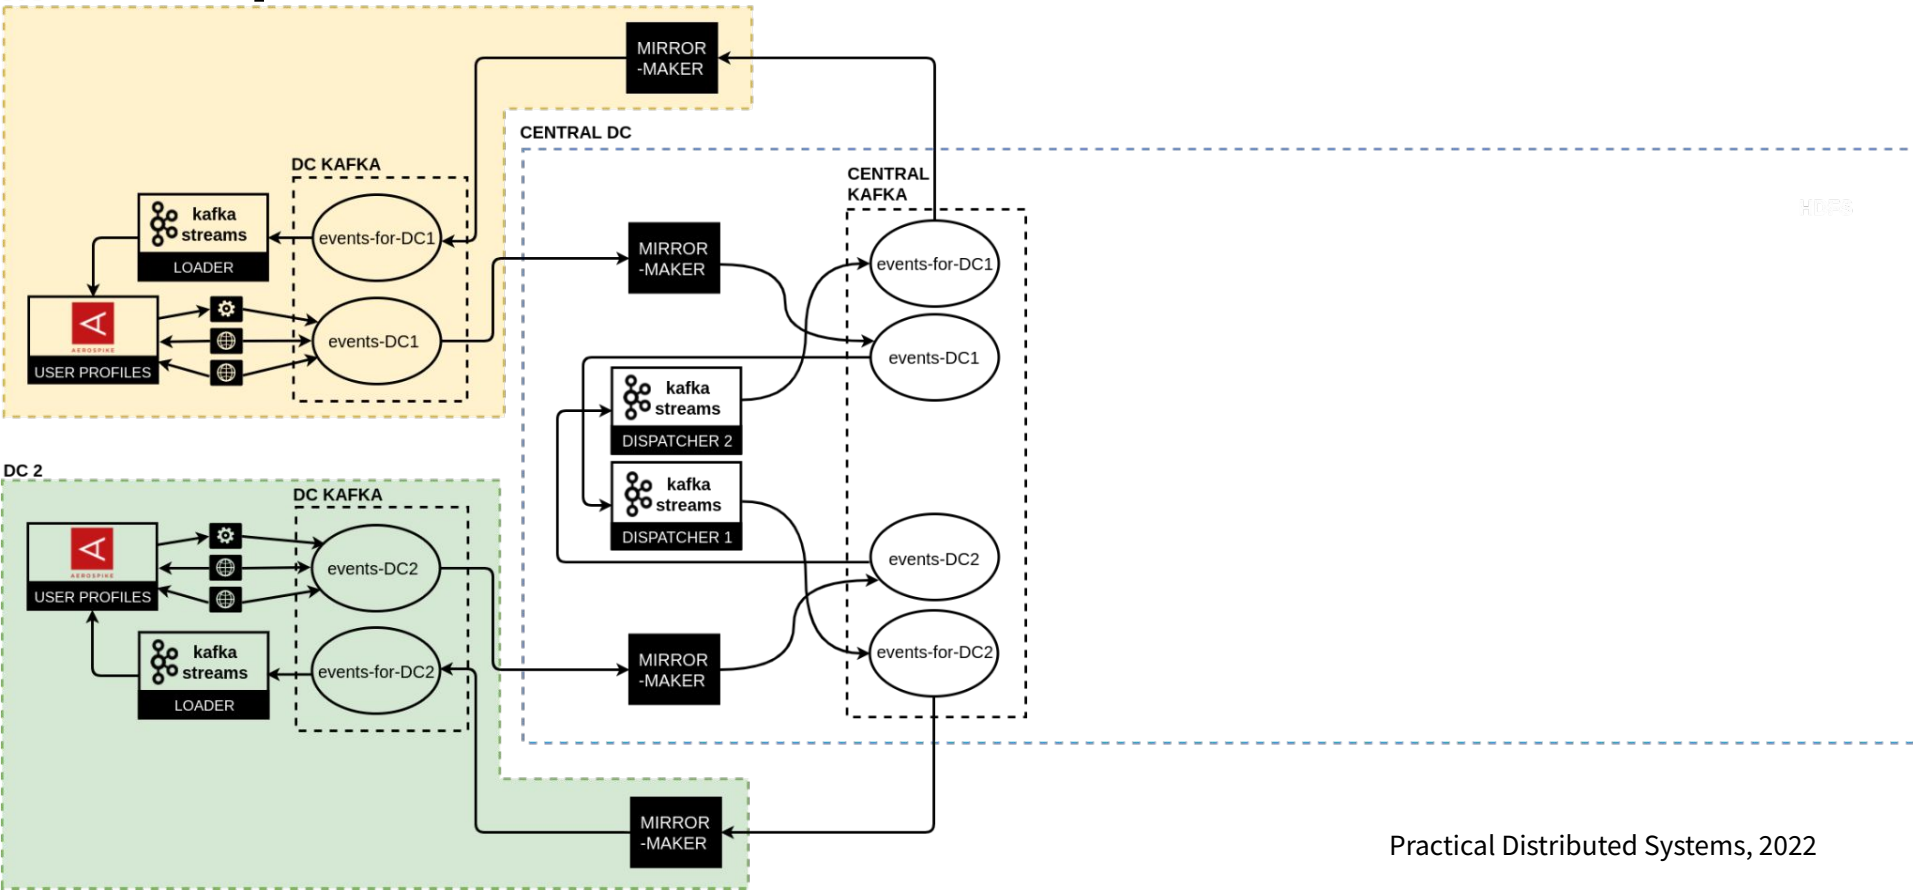


CENTRAL DC



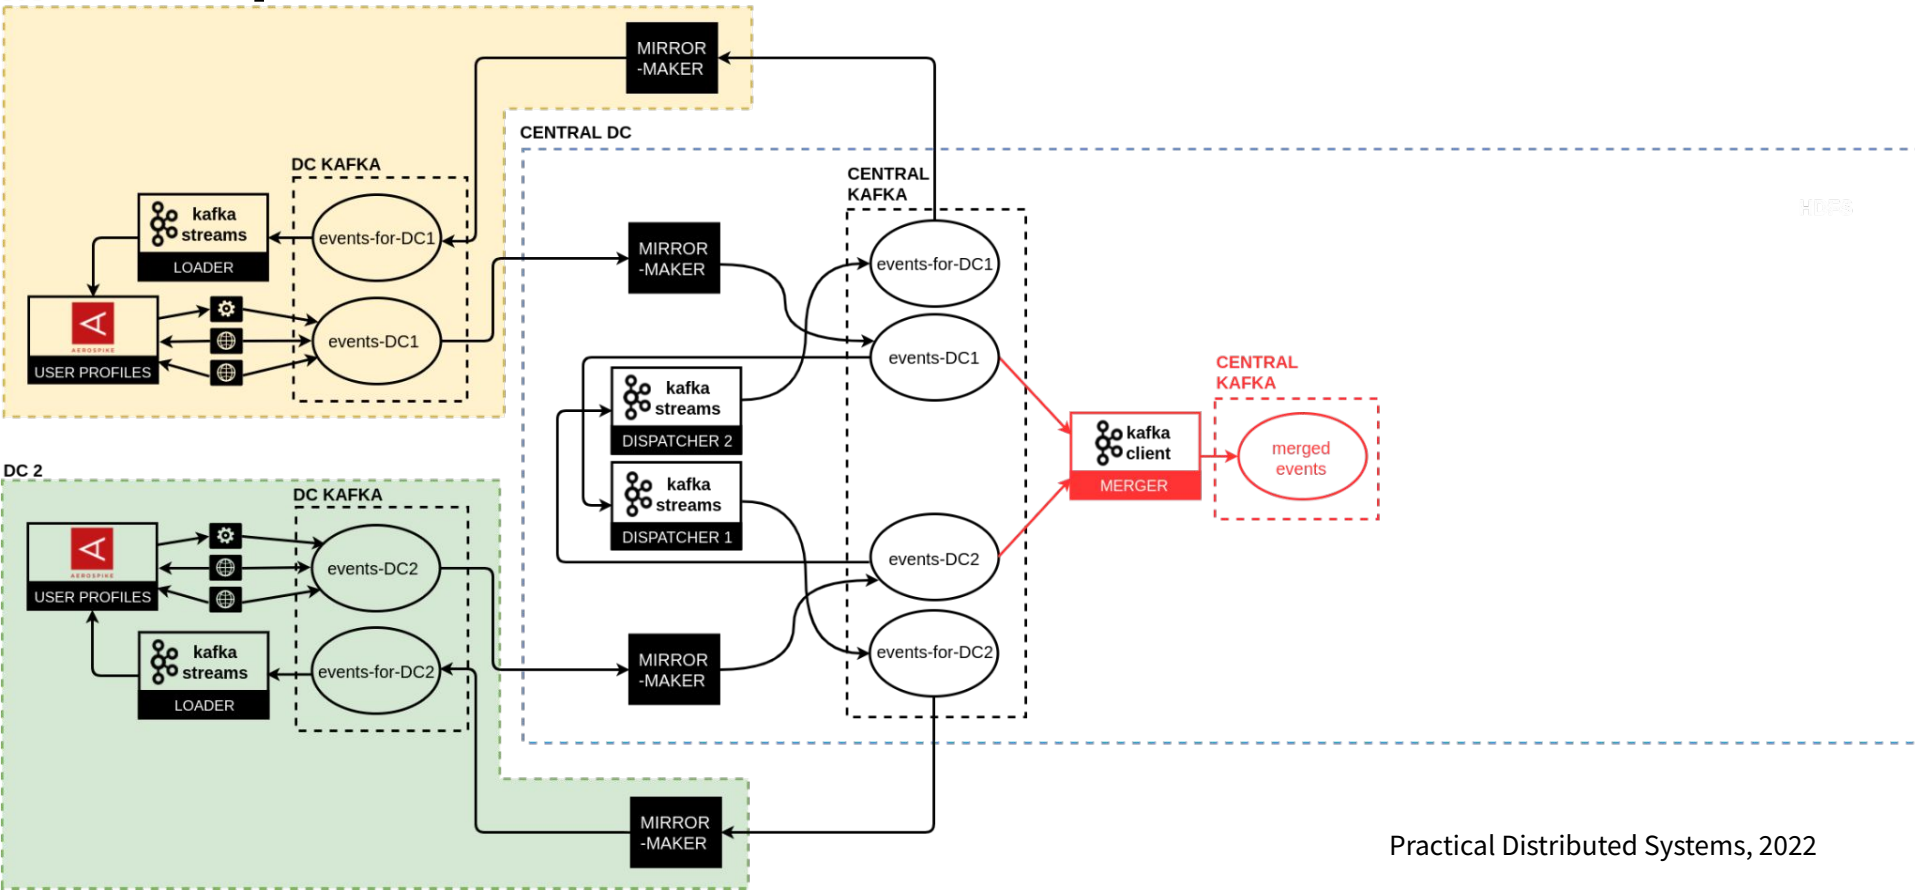
# The 4th iteration: multi-dc architecture

DC 1



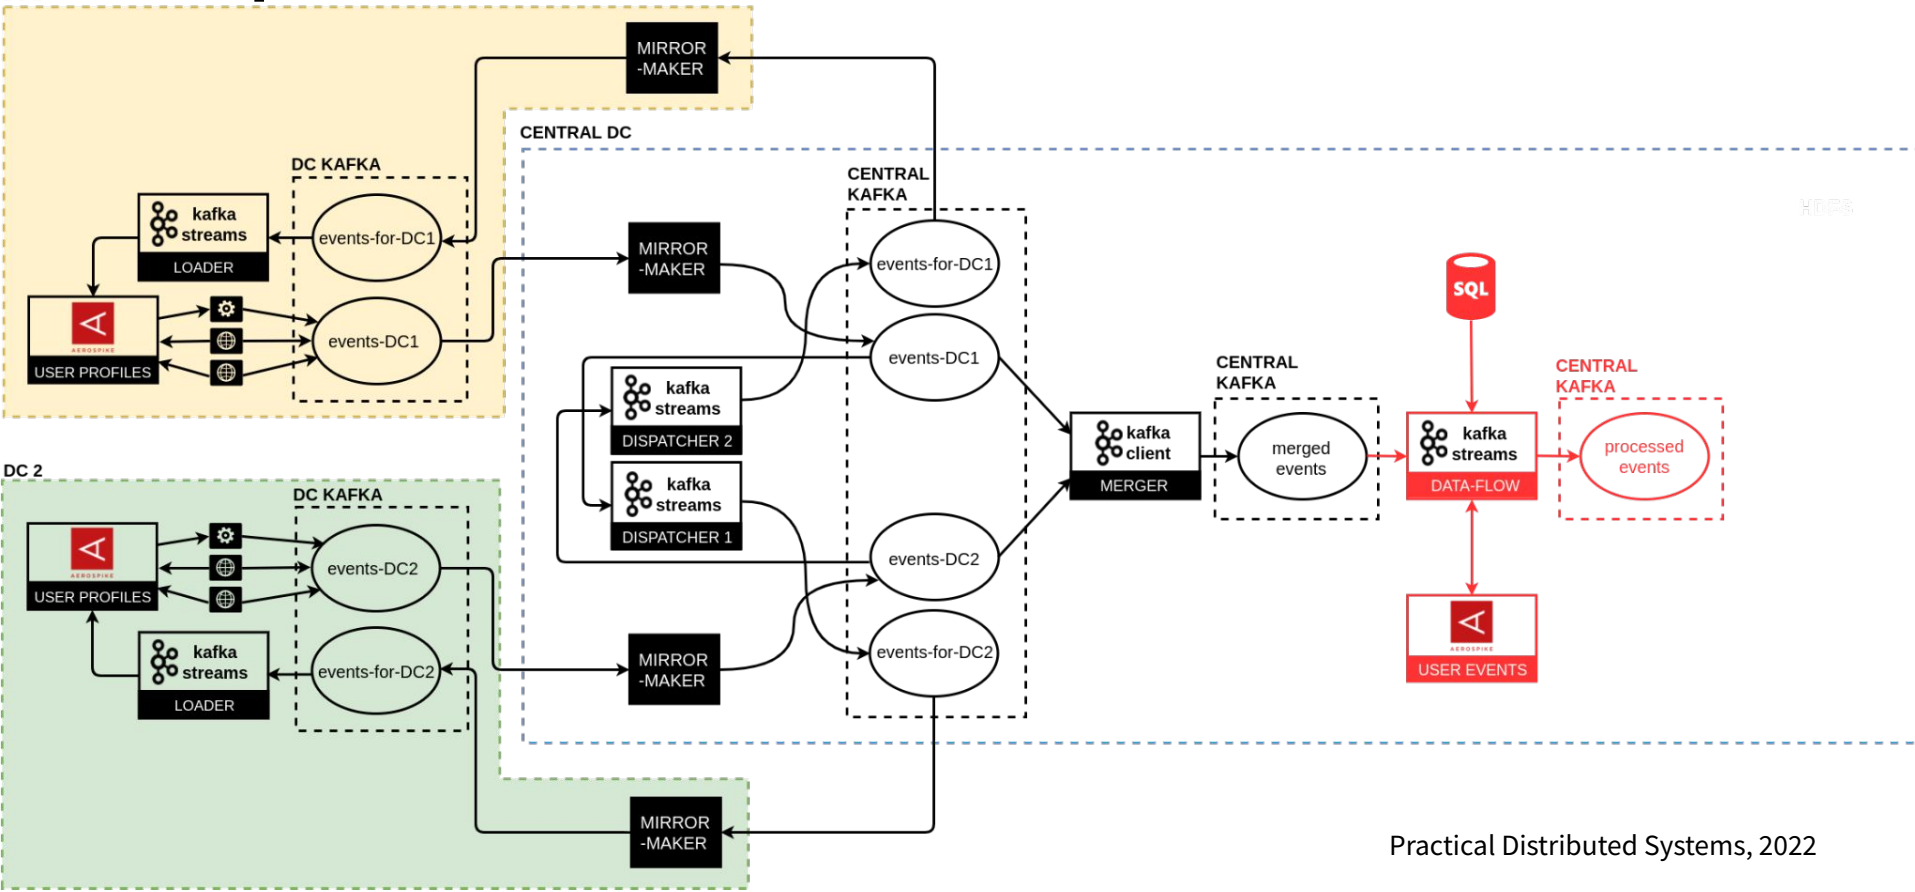
# The 4th iteration: multi-dc architecture

DC 1



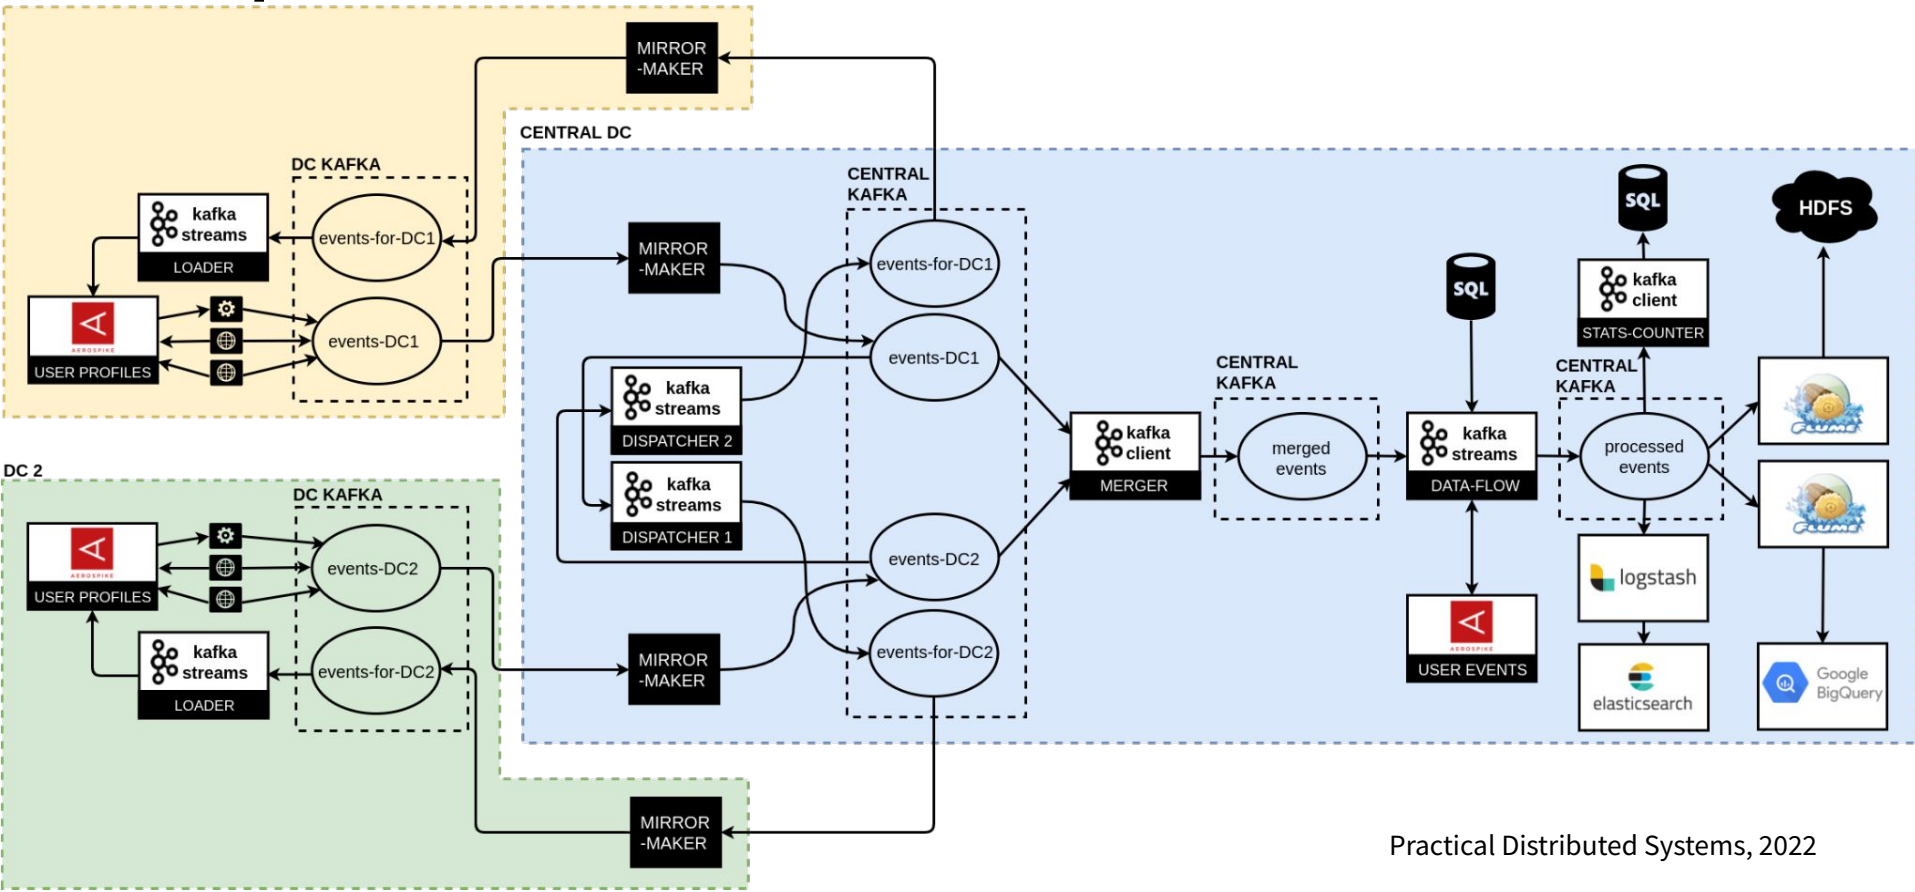
# The 4th iteration: multi-dc architecture

DC 1



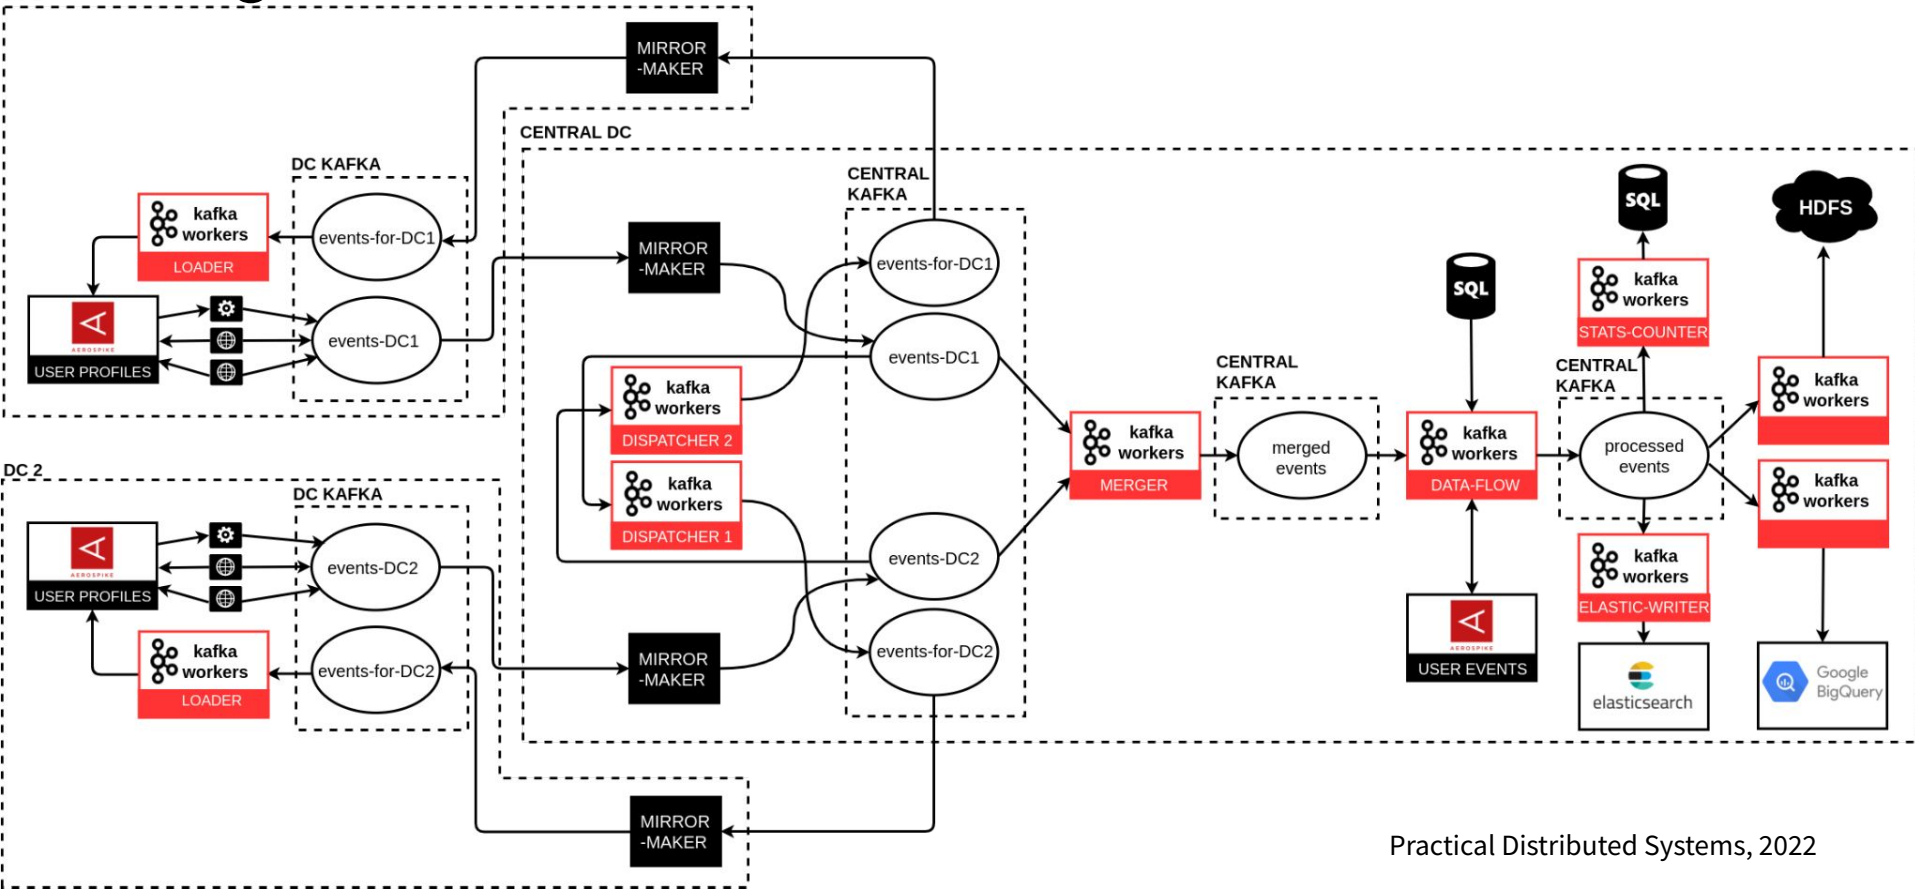
# The 4th iteration: multi-dc architecture

DC 1



# The 5th iteration: Kafka Workers


DC 1



# Apache Kafka

Why Kafka:

- event streaming platform (**distributed log**)
- producer-consumer separation
- fault-tolerance (**replication**)
- scalability and distribution (**topics partitioning**)
- log retention, statelessness
- efficient data consumption






# Apache Camus

Why Camus:

- **MapReduce** job that incrementally **loads data from Kafka into HDFS**
- fetches topics from Zookeeper and latest offsets from Kafka
- **partitions the output** based on the timestamp of each record
- **stores offsets** in log files in HDFS





# Apache Avro

Why Apache Avro:

- data serialization framework
- stores data in a compact, efficient **binary format**
- schema (JSON) could define rich data structures using various complex types
- schema is stored with data in one **Avro file** (self-describing container files)
- supports **schema changes** (old schema could be deserialized by a new program)

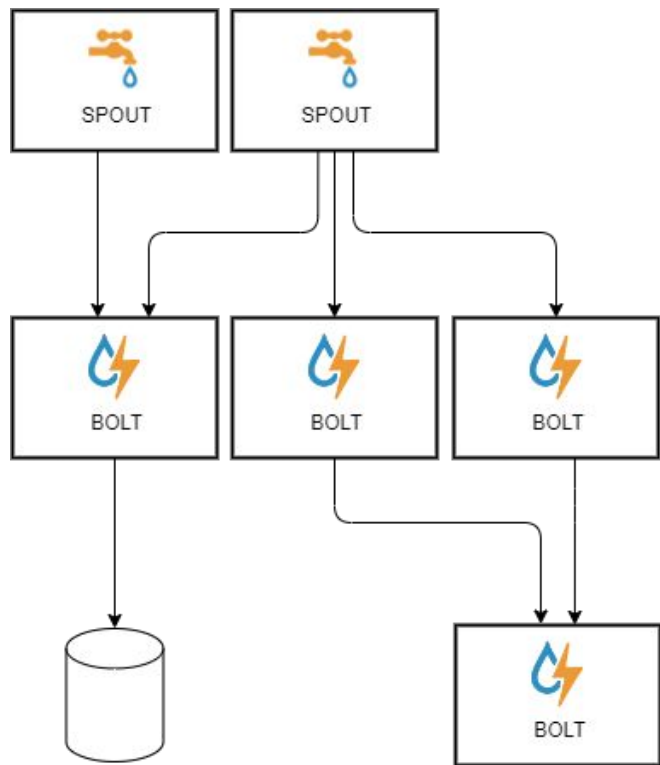
Our approach:

- Kafka's messages and HDFS files
- schema registry, historical schemas for Avro deserialization
- avro-fastserde ([github.com/RTBHOUSE/avro-fastserde](https://github.com/RTBHOUSE/avro-fastserde))

# Apache Storm

## Why Apache Storm:

- real-time computation system
- processes **streams** of **tuples** and runs user-defined **topologies** with processing nodes:
  - **spouts** emit new tuples
  - **bolts** receive tuples, do processing and generate tuples (**states** persist information)
- guarantees that every spout tuple will be fully processed (fault-tolerance)
- executes spouts and bolts as individual **tasks** that run in parallel on multiple machines



# Apache Storm

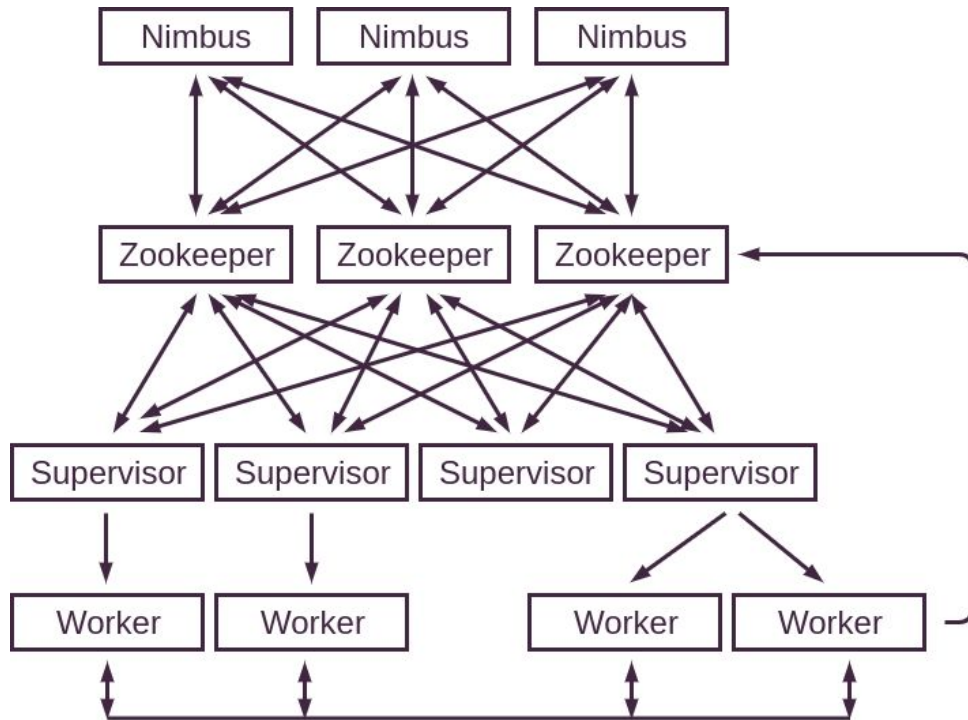
Master runs a daemon - **Nimbus**:

- responsible for distributing code around the cluster, assigning tasks to machines, and monitoring

Worker runs a daemon - **Supervisor**:

- listens for work assigned to its machine and starts and stops worker processes

Each **worker process** is a physical JVM and executes a subset of all the tasks for the topology



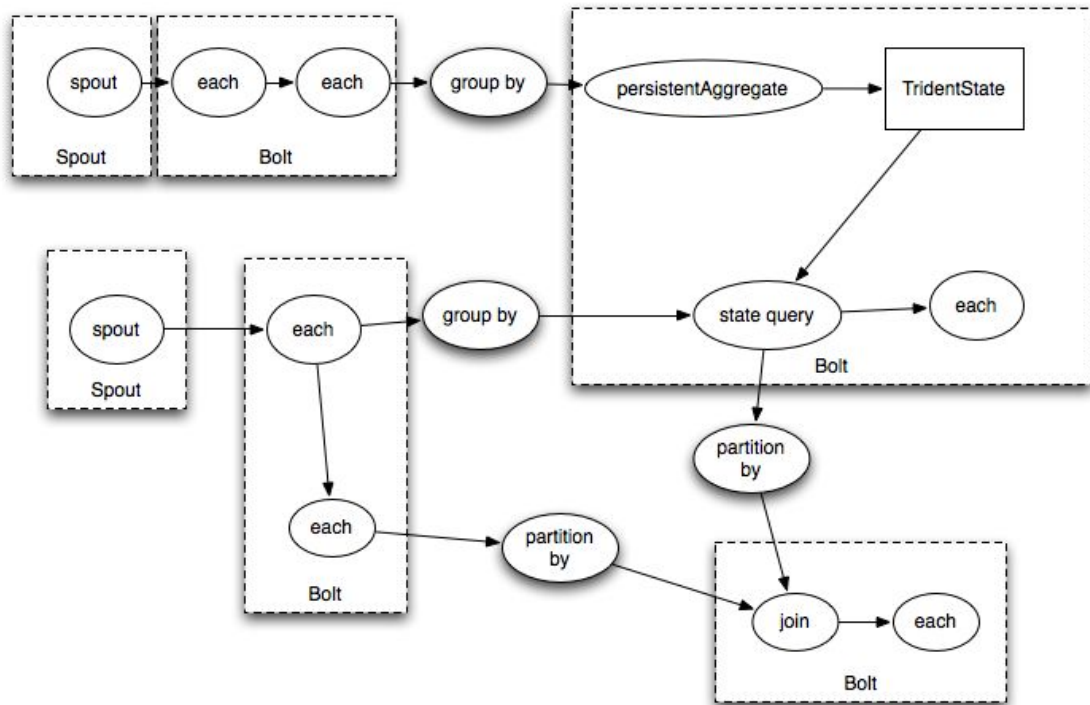


# Apache Storm: Trident API

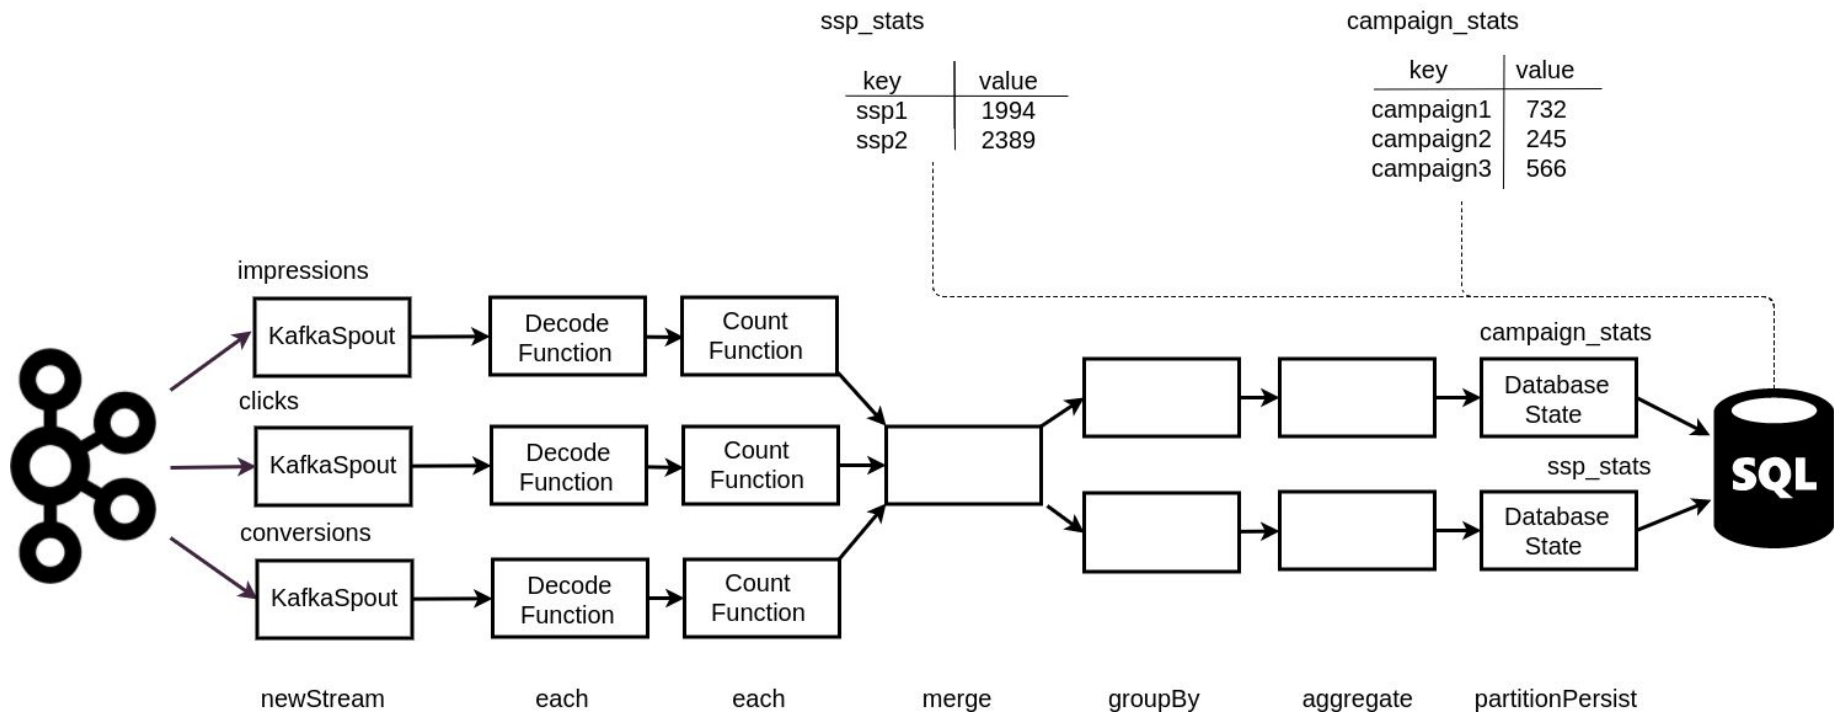
Source: storm.apache.org

## Why Trident:

- high-level declarative API
- provides **functions, filters, joins, groupings, and aggregations**
- supports stateful, incremental processing on top of persistence stores
- processes **microbatches** (transactions) and supports exactly-once processing



# Use case: stats-counter



# Use case: stats-counter (exactly-once state)

zookeeper

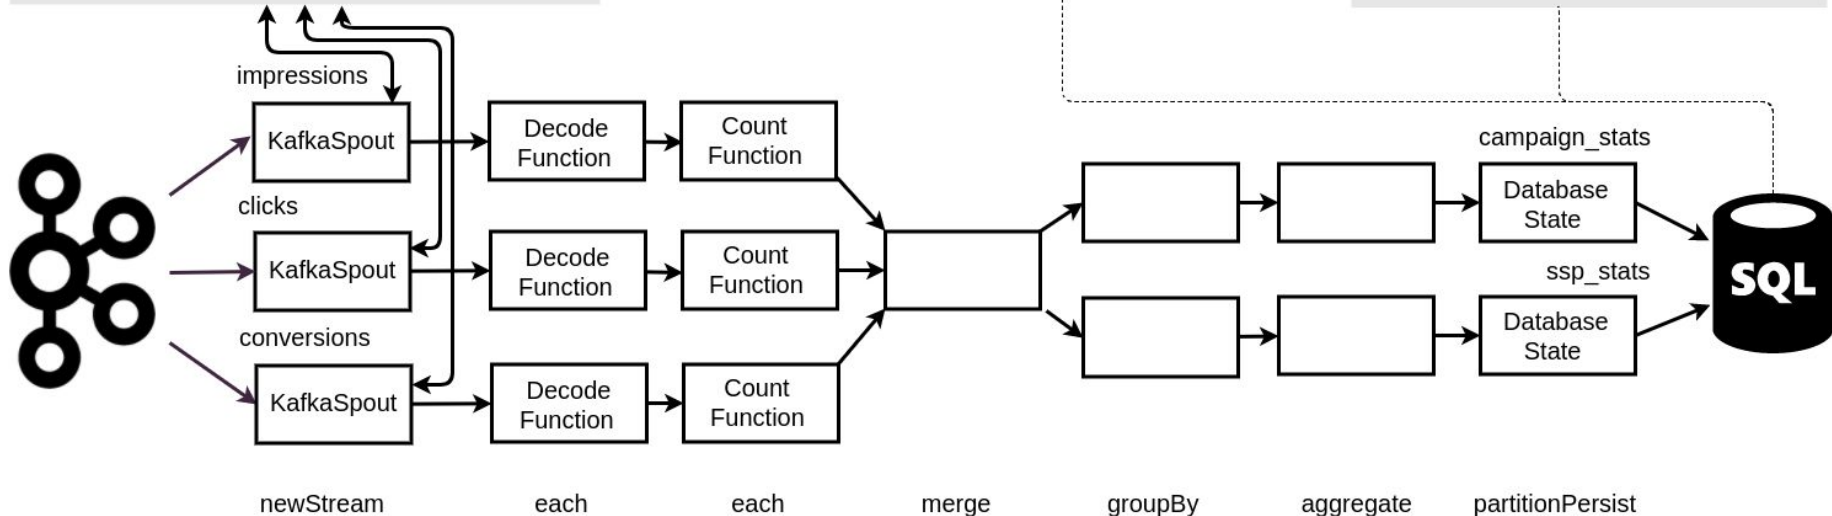
topic	<offset1, offset2>	transaction_id
impressions	<5831, 5917>	19
clicks	<623, 680>	19
conversions	<423, 442>	19

ssp\_stats

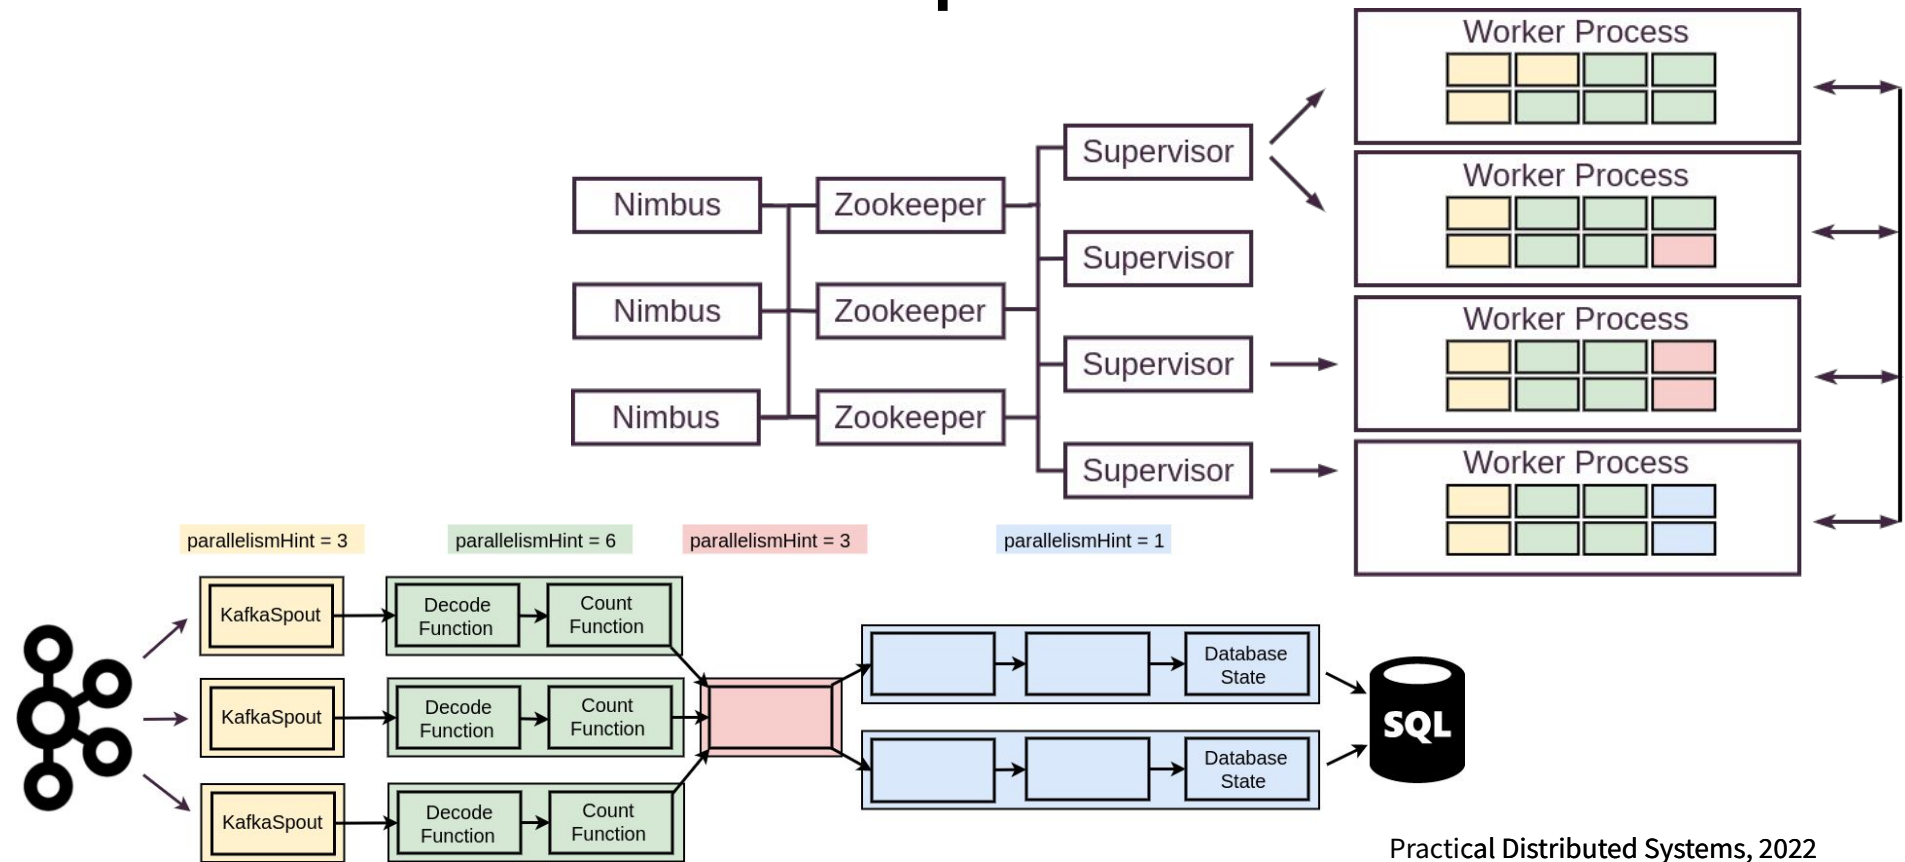
key	value	last_transaction_id
ssp1	1994	17
ssp2	2389	18

campaign\_stats

key	value	last_transaction_id
campaign1	732	18
campaign2	245	17
campaign3	566	18



# Use case: stats-counter (parallelism)






# Kafka Streams

Source: kafka.apache.org

## Why Kafka Streams:

- Java library (based on Kafka **producer and consumer APIs**) run as a standard application
- **no processing cluster** and no external dependencies
- uses Kafka's **parallelism model** and group membership mechanism (scalability and fault-tolerance)
- does **event-at-a-time** processing (no batching)
- supports **exactly-once** processing



# Kafka Streams: topology

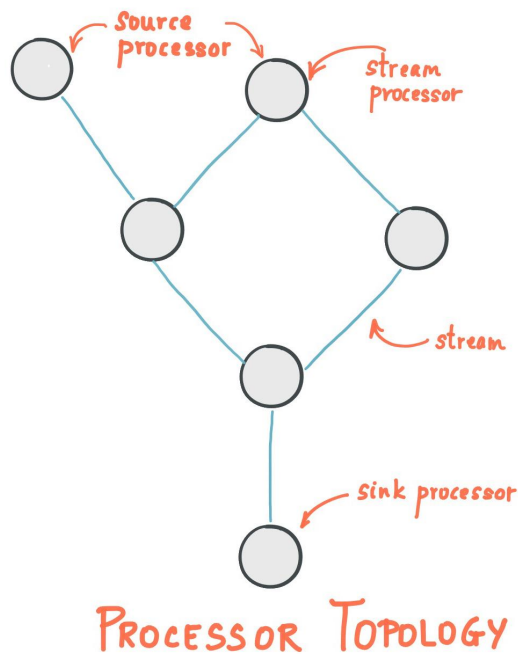
Source: kafka.apache.org

Topology is a graph of **stream processors** that are connected by **streams**:

- consumes records from one or more input Kafka topics (source processors)
- sends records to output Kafka topic (sink processors)

Kafka Streams uses Kafka concepts:



- **data record** - Kafka message
- **stream partition** - Kafka topic partition
- **keys** determine the partitioning



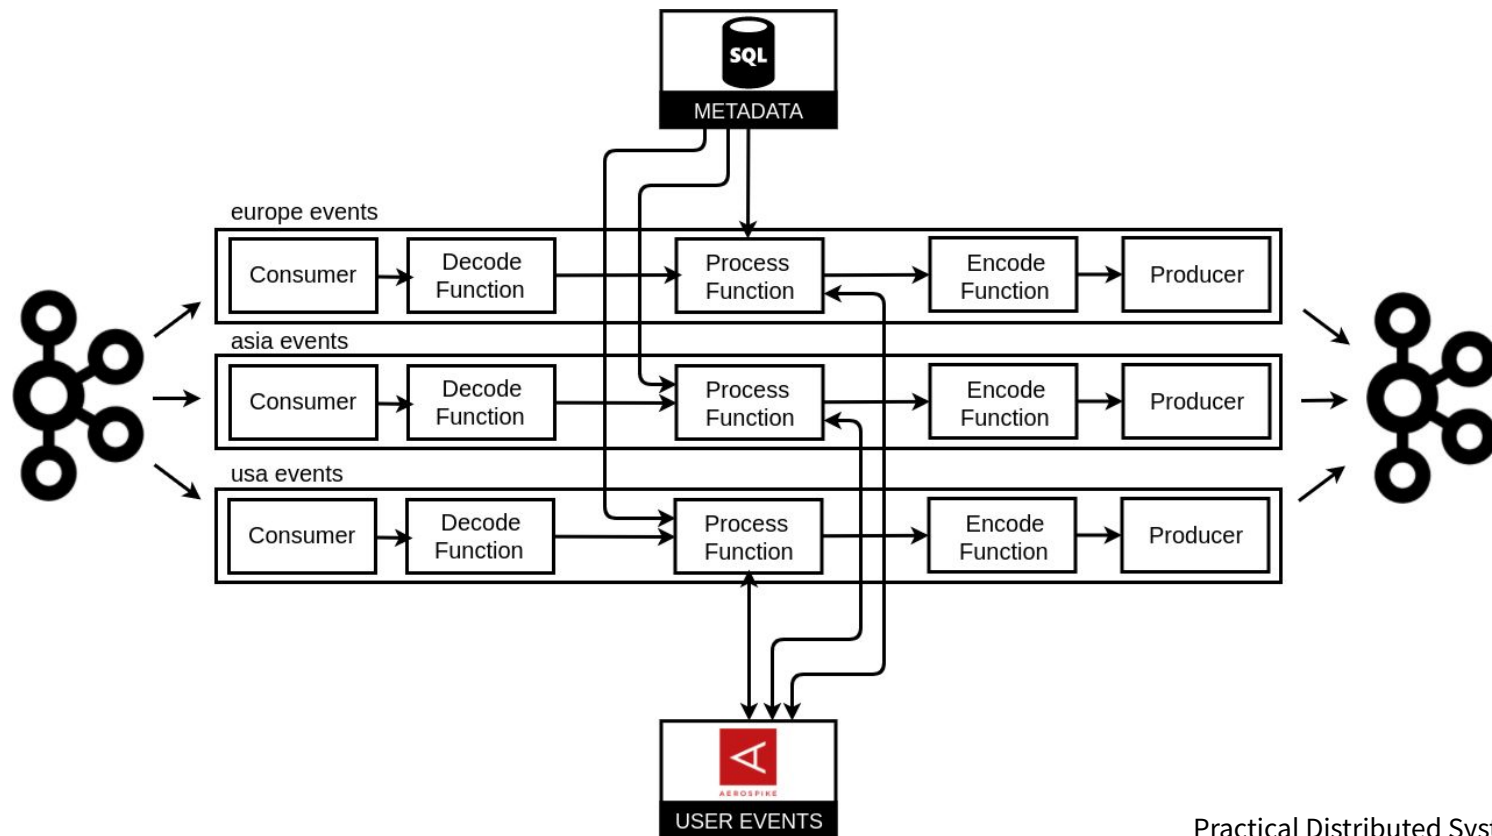


# Kafka Streams: threading model


Source: [kafka.apache.org](https://kafka.apache.org)



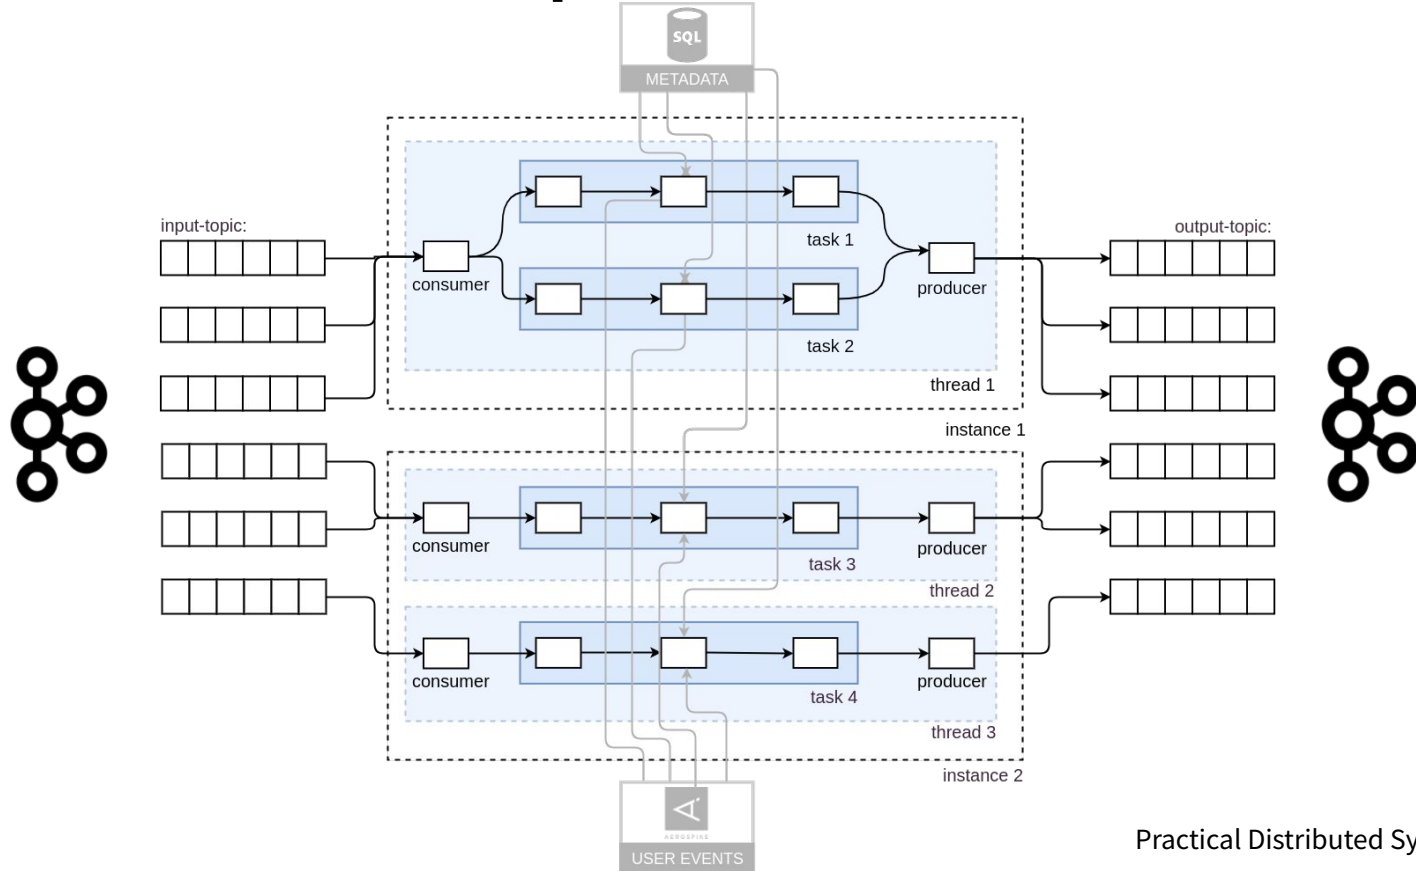
# Use case: data-flow




# Use case: data-flow (parallelism)

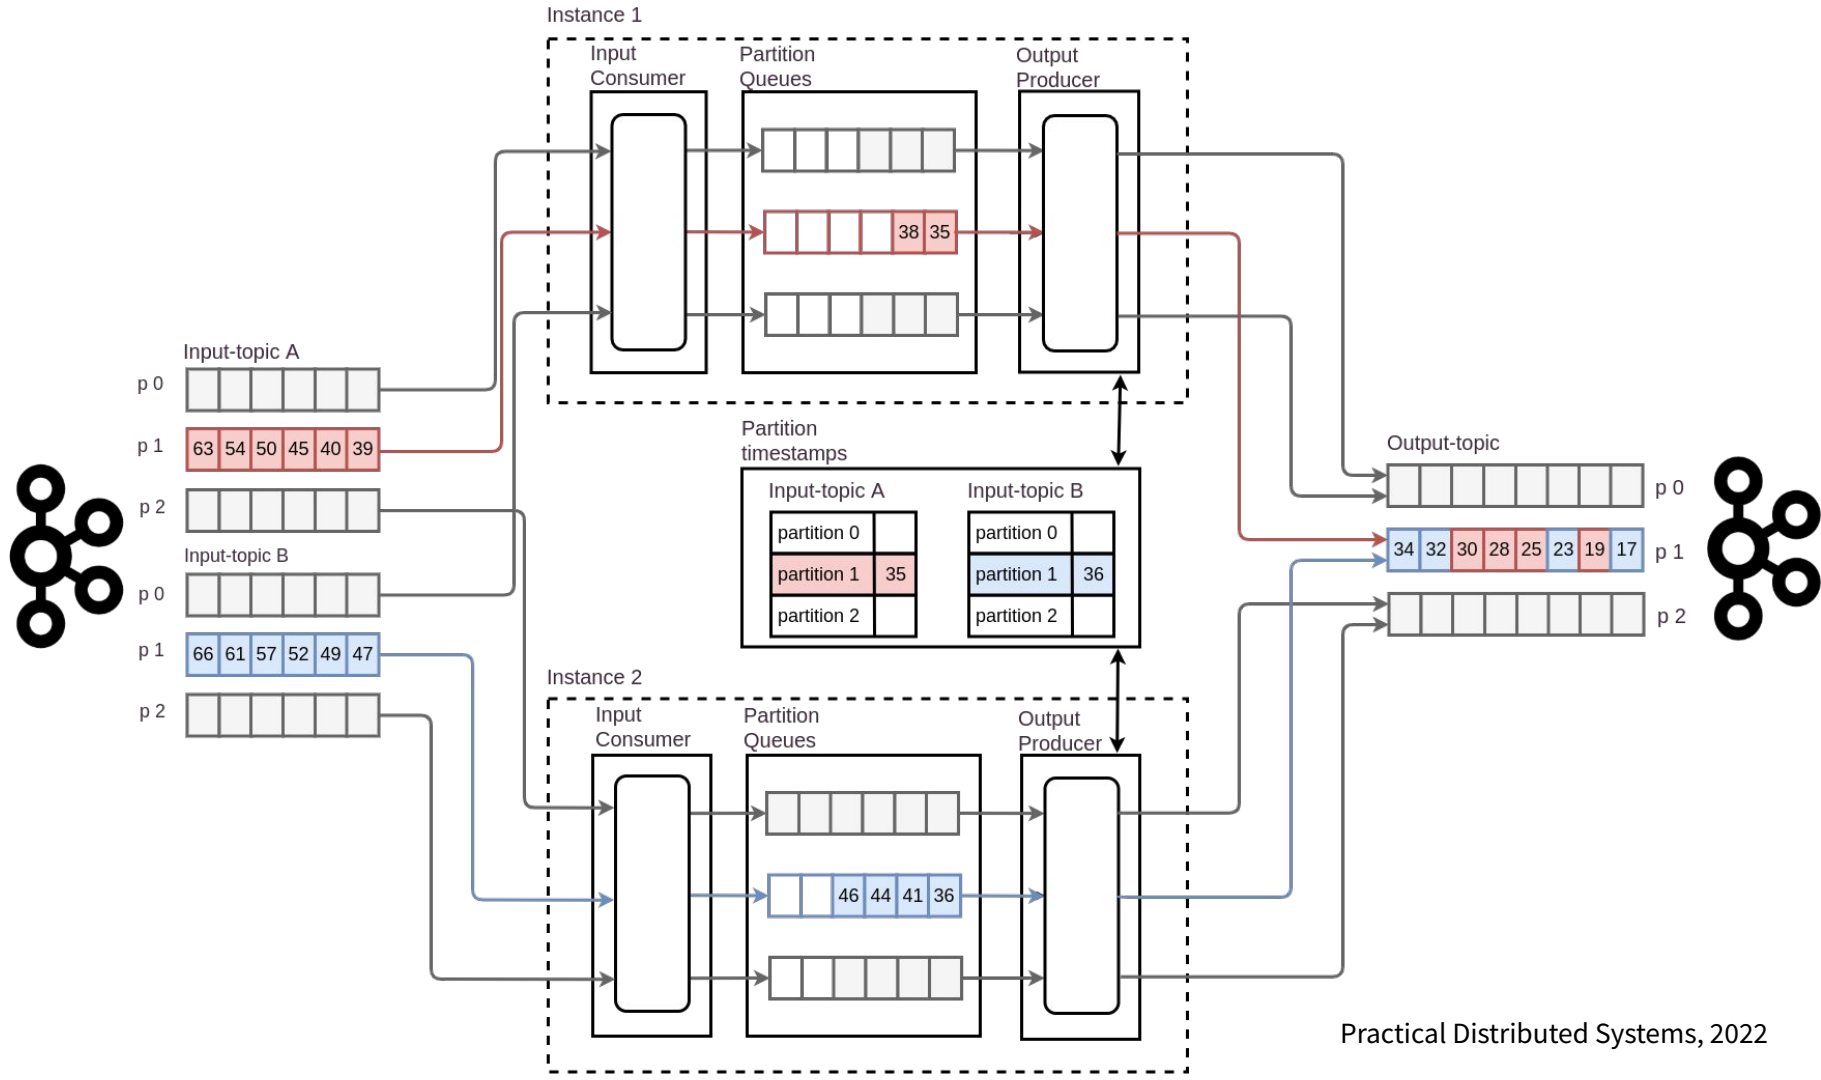


# Use case: data-flow (parallelism)



# Use case: merger



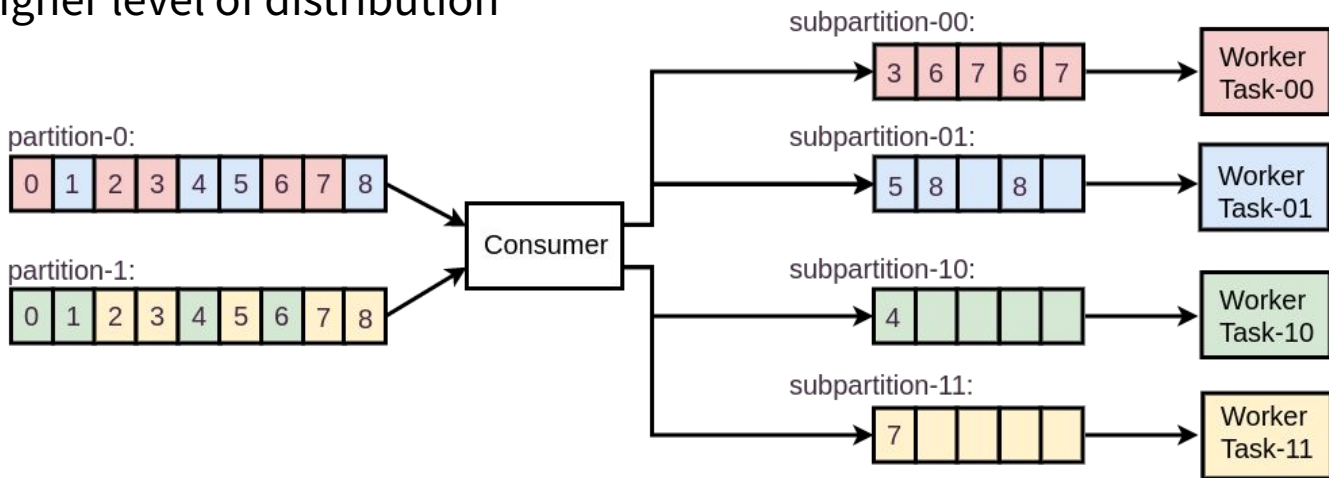




# Kafka Workers: main features

Why Kafka Workers ([github.com/RTBHOUSE/kafka-workers](https://github.com/RTBHOUSE/kafka-workers)):


- better threading model with better resources utilization
  - separating processing from consumption
  - higher level of distribution



# Kafka Workers: main features

Why Kafka Workers ([github.com/RTBHOUSE/kafka-workers](https://github.com/RTBHOUSE/kafka-workers)):

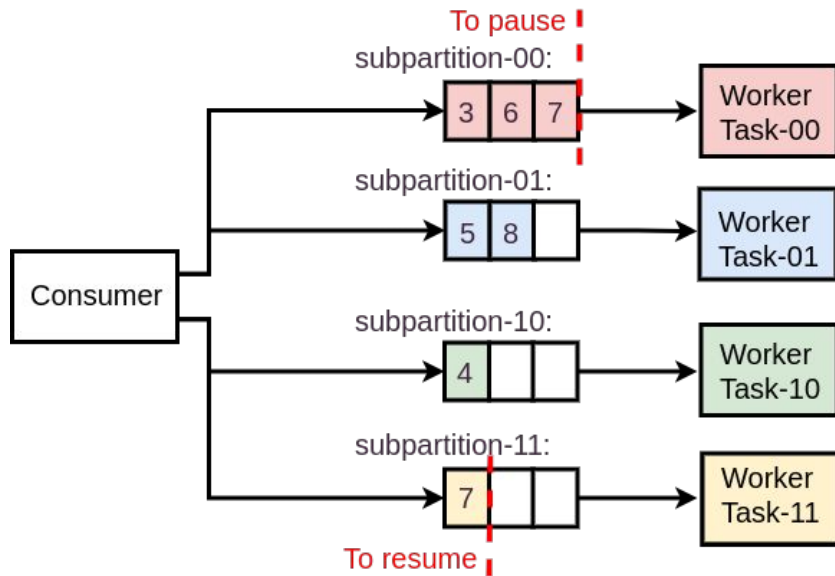
- asynchronous processing
  - processing timeouts
  - tighter control of offset commits



# Kafka Workers: main features

Why Kafka Workers ([github.com/RTBHOUSE/kafka-workers](https://github.com/RTBHOUSE/kafka-workers)):

- backpressure



# Kafka Workers: main features

Why Kafka Workers ([github.com/RTBHOUSE/kafka-workers](https://github.com/RTBHOUSE/kafka-workers)):

- possibility to pause and resume processing for a given partition
- at-least-once semantics
  - handling failures
- simplicity
  - Kafka Consumer API
  - no processing cluster, no external dependencies
  - without translating messages to/from its internal data format
  - no interprocess communication
- kafka-to-kafka, hdfs, bigquery, elasticsearch connectors

# Kafka Workers: API (subpartitions)

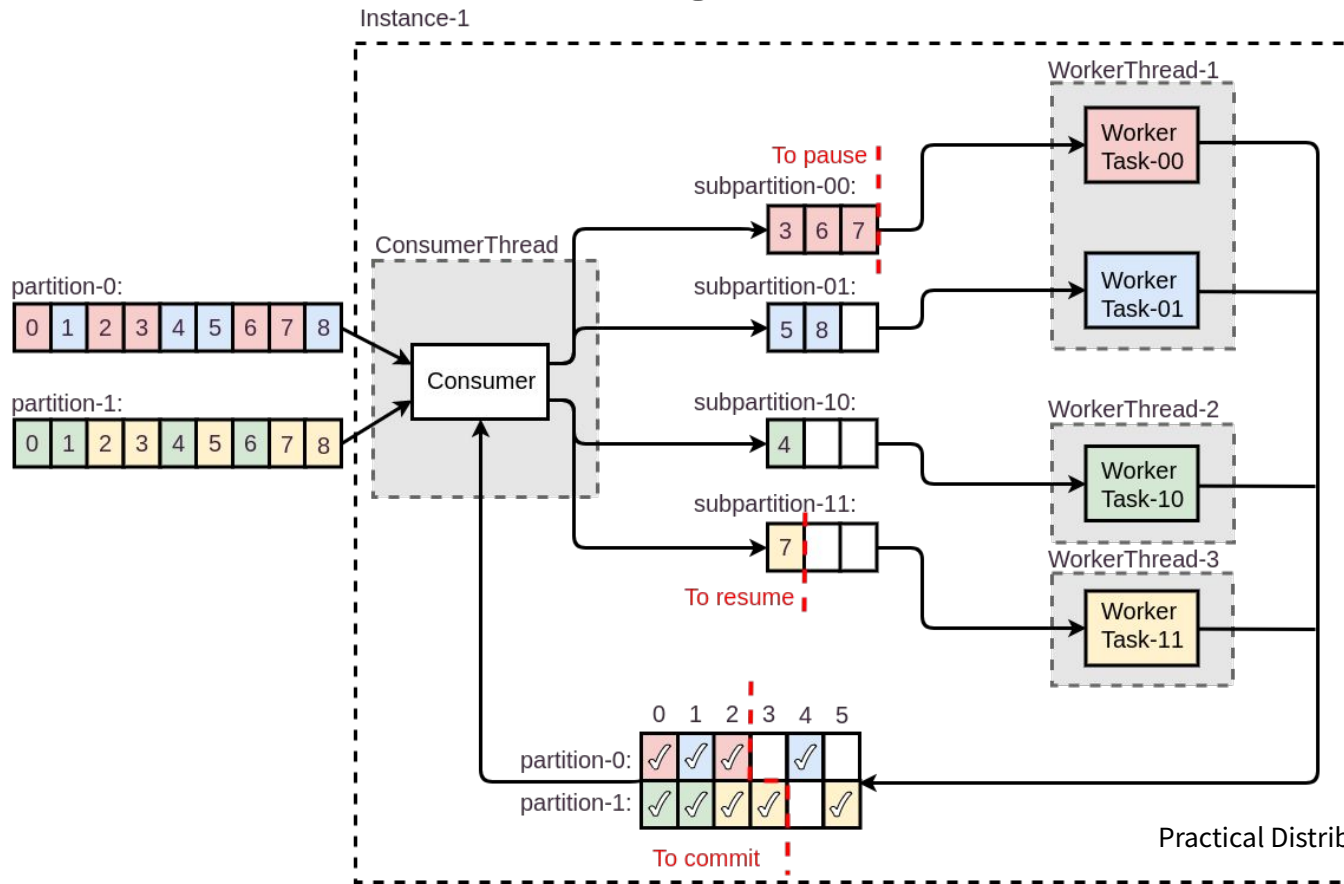
```
public interface WorkerPartitioner<K, V> {  
    int subpartition(ConsumerRecord<K, V> consumerRecord);  
}
```

# Kafka Workers: API (tasks)

```
public interface WorkerTask<K, V> {  
    boolean accept(WorkerRecord<K, V> record);  
    void process(WorkerRecord<K, V> record, RecordStatusObserver observer);  
}
```

```
public interface RecordStatusObserver {  
    void onSuccess();  
    void onFailure(Exception exception);  
}
```

# Kafka Workers: threading model



# Summary

What we have achieved:

- platform monitoring
- much more stable platform
- higher quality of data processing
- HDFS & BigQuery & Elasticsearch streaming
- multi-DC architecture and data synchronization
- high scalability
- better data-flow monitoring, deployment & maintenance

